

Parallel STL – summary

Execution policies

- `sequenced_policy` (`std::execution::seq`): execution may not be parallelized
- `parallel_policy` (`std::execution::par`): execution may be parallelized
- `parallel_unsequenced_policy` (`std::execution::par_unseq`): execution may be parallelized, element access functions may be invoked in any order, the execution may be vectorized
- `unsequenced_policy` (`std::execution::unseq`, since C++20): execution may be vectorized

Performance summary

The operation used for this example is a `std::transform` with a relatively complex task executed on a `std::vector` of integers. Results show the number of iterations per second.

Data size	x86 seq	x86 par	arm64 seq	arm64 par
32	354k (100%)	74.0k (21%)	1.81M (100%)	242k (13%)
128	66.3k (100%)	33.2k (50%)	436k (100%)	86.9k (20%)
512	18.7k (100%)	16.7k (89%)	96.5k (100%)	45.1k (47%)
2048	4.70k (100%)	8.12k (173%)	17.7k (100%)	32.9k (186%)
8192	1.10k (100%)	2.92k (265%)	4.20k (100%)	14.7k (350%)
32768	277 (100%)	835 (301%)	993 (100%)	4.53k (456%)
131072	65.5 (100%)	222 (339%)	242 (100%)	1.15k (475%)

Compilers and HW used for testing: x86-64 – MSVC 2019 (/O2) on Intel i7-8650U (4 cores × 2 threads), ARM64 – GCC 10.2 (-O3) on Apple M1 (4 low power + 4 high power cores).

Basic observations

- Data size and platform should be taken into consideration when deciding on an execution policy. Spawning threads itself has a cost, which may be different between platforms.
- Manually setting the `unseq` policy seemed to provide very little (if any) profit. This is due to the compiler itself optimizing code, when `-O0` option is used to disable optimizations, the gain from this policy is noticeably higher.